




# KubeDay

JAPAN





**KubeDay**  
JAPAN

# **Secure and Debuggable: Debugging Slim, Scratch and Distroless Kubernetes Containers**

*Saiyam Pathak & Kyle Quest*

# Intro

- **Expected** background
  - basic container fundamentals
  - basic knowledge of namespaces
  - k8s fundamental
- **What** will be covered
- **Who** we are

# About me



- Director of Technical Evangelism, Civo
- CNCF Ambassador
- CKA, CKAD, CKS Certified
- KCNA Certification SME

 [saiyampathak.com/twitter](https://saiyampathak.com/twitter)


 [saiyampathak.com/youtube](https://saiyampathak.com/youtube)

 [saiyampathak.com/discord](https://saiyampathak.com/discord)

 [kubesimplify.com](https://kubesimplify.com)

## Let's Learn CKS Scenarios

Author of a book which  
helps you prepare for the  
Kubernetes CKS exam





## Kyle Quest

(aka the docker-slim guy)

- DockerSlim creator
- Go enthusiast
  - 50 Shades of Go
  - Go Project Layout
- eBPF hacker
- Slim.AI founder/CTO
- <https://twitter.com/kcqon>
- <https://github.com/kcq>

# The Problem

“

This minimal container  
image worked fine, but  
only if I never made a  
mistake.



# General Debugging Techniques for Kubernetes

- Finding/tracking failures in kubernetes (events, logs, kubectl commands)
- Host/node level debugging
  - traditional (ssh-based)
  - using privileged containers and nsenter
- Embedding debugging tools
- Copying debugging tools
- Interactive container debugging with kubectl exec
- Permanent observability/debugging sidecars

## Troubleshooting containers

STATUS	RESTARTS
Pending	0
Running	0
CrashLoopBackOff	1 (7s ago)

crashLoopBackOff

kubectl describe

kubectl logs pod --previous

kubctl exec -it pod -- sh/bash

Kubectl get events

See events  
Pending pod

Insufficient resources

For running containers  
with debugging utilities

```
Events:
  Type      Reason      Age   From              Message
  ----      -
Warning    FailedScheduling 21s   default-scheduler 0/3 nodes are available: 3 node(s) didn't match Pod's node affinity/selector
```

```
demo ~ kubectl exec -it nginx -- sh
# ls
bin  docker-entrypoint.d  home  media  proc  sbin  tmp
boot docker-entrypoint.sh lib   mnt    root  srv   usr
dev  etc                  lib64 opt    run   sys   var
```

# Slim and Minimal Container Image

- **Why**
- **How**/options
  - Scratch
  - Distroless
  - Slim
- **Gotchas**

# Scratch

Empty image  
Nothing in it

Mostly used in multi-stage build

# Distroless

Not completely empty  
/etc/passwd  
CA certs

only app  
+  
runtime  
dependencies

No shell  
No Package manger

## More Gotchas



→ Extra dependencies(complicated)

→ No debugging tooling

→ No Shell

→ Not always static binary

# Creating Minimal Container Images



```
% nerdctl -n k8s.io images
```

REPOSITORY	TAG	IMAGE ID	SIZE
ghcr.io/kubeday-japan/demo-node-app	fat	13db159790c7	1.0 GiB
ghcr.io/kubeday-japan/demo-node-app	distroless	b9ce468dce45	166.8 MiB
ghcr.io/kubeday-japan/demo-node-app	slim	412eeb9bdb61	90.9 MiB



File Changes for All 1 Layers						
Search		1,967 Objects	Columns	TREE	FLAT	
↑ Name	Added	Modified	Deleted	Count	Mode	Size
▼ /	0	-	-	1,848	drwxrwxrwx	2.3 MB
▶ .	0	-	-	0	drwxr-xr-x	0
▶ bin	0	-	-	0	drwxr-xr-x	0
▶ boot	0	-	-	0	drwxr-xr-x	0
▶ dev	0	-	-	0	drwxr-xr-x	0
▶ etc	0	-	-	15	drwxr-xr-x	219.6 kB
▶ home	0	-	-	0	drwxr-xr-x	0
▶ lib	0	-	-	0	drwxr-xr-x	0
▶ proc	0	-	-	0	drwxr-xr-x	0
▶ root	0	-	-	0	drwx-----	0
▶ run	0	-	-	0	drwxr-xr-x	0
▶ sbin	0	-	-	0	drwxr-xr-x	0
▶ sys	0	-	-	0	drwxr-xr-x	0
▶ tmp	0	-	-	0	drwxrwxrwx	0
▶ usr	0	-	-	1,830	drwxr-xr-x	2.1 MB
▶ var	0	-	-	3	drwxr-xr-x	1.8 kB

**distroless/static** - 2.34MB

- Basic FS layout
- **/usr/share/zoneinfo** - 1.7 MB

File Changes for All 2 Layers						
<div> <div>Search</div> <div>2,443 Objects</div> <div>Columns</div> <div>TREE</div> <div>FLAT</div> </div>						
↑ Name	Added	Modified	Deleted	Count	Mode	Size
▼ /	0	-	-	2,307	drwxrwxrwx	20.3 MB
▶ .	0	1	-	0	drwxr-xr-x	0
▶ bin	0	-	-	0	drwxr-xr-x	0
▶ boot	0	-	-	0	drwxr-xr-x	0
▶ dev	0	-	-	0	drwxr-xr-x	0
▶ etc	0	1	-	17	drwxr-xr-x	230.9 kB
▶ home	0	-	-	0	drwxr-xr-x	0
▶ lib	0	1	-	37	drwxr-xr-x	4.3 MB
▶ lib64	1	-	-	1	drwxr-xr-x	0
▶ proc	0	-	-	0	drwxr-xr-x	0
▶ root	0	-	-	0	drwx-----	0
▶ run	0	-	-	0	drwxr-xr-x	0
▶ sbin	0	-	-	0	drwxr-xr-x	0
▶ sys	0	-	-	0	drwxr-xr-x	0
▶ tmp	0	-	-	0	drwxrwxrwx	0
▶ usr	0	1	-	2,246	drwxr-xr-x	15.8 MB
▶ var	0	1	-	6	drwxr-xr-x	4.7 kB

**distroless/base** - 20.3MB

- **distroless/static++**
- **/lib/x86\_64-linux-gnu/** - 4.3MB
- **/usr/lib/x86\_64-linux-gnu/** - 11.7MB

File Changes for All 4 Layers						
Search		3,444 Objects		Columns	TREE	FLAT
↑ Name	Added	Modified	Deleted	Count	Mode	Size
▼ /	0	-	-	2,862	drwxrwxrwx	110.0 MB
▶ .	0	1, 2, 3	-	0	drwxr-xr-x	0
▶ bin	0	-	-	0	drwxr-xr-x	0
▶ boot	0	-	-	0	drwxr-xr-x	0
▶ dev	0	-	-	0	drwxr-xr-x	0
▶ etc	0	1	-	17	drwxr-xr-x	230.9 kB
▶ home	0	-	-	0	drwxr-xr-x	0
▶ lib	0	1, 2	-	38	drwxr-xr-x	4.4 MB
▶ lib64	1	-	-	1	drwxr-xr-x	0
▶ nodejs	3	-	-	538	drwxr-xr-x	87.4 MB
▶ proc	0	-	-	0	drwxr-xr-x	0
▶ root	0	-	-	0	drwx-----	0
▶ run	0	-	-	0	drwxr-xr-x	0
▶ sbin	0	-	-	0	drwxr-xr-x	0
▶ sys	0	-	-	0	drwxr-xr-x	0
▶ tmp	0	-	-	0	drwxrwxrwx	0
▶ usr	0	1, 2	-	2,259	drwxr-xr-x	18.0 MB
▶ var	0	1, 2	-	9	drwxr-xr-x	6.6 kB

**distroless/nodejs - 110MB**

- distroless/base++
- /usr/lib/x86\_64-linux-gnu/ - 13.9MB
- /nodejs/ - 87.4MB



```
FROM node:18 as builder
```

```
WORKDIR /usr/src/app
```

```
COPY app/package*.json ./
```

```
RUN npm install
```

```
COPY app/ .
```

```
FROM gcr.io/distroless/nodejs:18
```

```
WORKDIR /usr/src/app
```

```
COPY --from=builder /usr/src/app .
```

```
EXPOSE 8080
```

```
CMD [ "node", "server.js" ]
```

## **distroless app** - 166.8MB

- distroless/nodejs++
- COPY “everything”

# Create Minimal Container Images with Slim (the easy way)



kubeday-japan / mux-go-api-ubuntu

latest @ sha256:c08720...f699b5



## Harden the image



kubeday-japan / mux-go-api-ubuntu

latest-slim

From Slim Cloud

Overview



File Explorer



Dockerfile



Vulnerabilities



Layer All Layer 0: 63.1 MB Layer 1: 559.5 MB Layer 2: 0 B Layer 3: 3.4 kB Layer 4: 14.9 MB

Instructions for All 5 Layers

Container Startup Files

[/root/go/src/api/api](#)  
[/root/go/src/api](#)

File Changes for All 5 Layers

Search

19,791 Objects

Columns

TREE

FLAT

↑ Name	Added	Modified	Deleted	Count	Mode	Size
▼ /	0	-	-	17,001 drwxrwxrwx		636.8 MB
▶ bin	0	1	-	90 drwxr-xr-x		5.1 MB
▶ boot	0	-	-	0 drwxr-xr-x		0
▶ dev	0	-	-	0 drwxr-xr-x		0
▶ etc	0	1	-	423 drwxr-xr-x		933.9 kB
▶ home	0	-	-	0 drwxr-xr-x		0
▶ lib	0	1	-	220 drwxr-xr-x		12.8 MB

## Output summary

Your container was successfully optimized. The new image is **98%** smaller than the source image.

Container size before: **638 MB**

Container size after: **9.9 MB**

Minification: **64.56x**

DOWNLOAD

File Changes for All 1 Layers						
<div> <div>Search</div> <div>424 Objects</div> <div>Columns</div> <div>TREE</div> <div>FLAT</div> </div>						
↑ Name	Added	Modified	Deleted	Count	Mode	Size
▼ /	0	-	-	403	drwxrwxrwx	9.9 MB
▶ etc	0	-	-	259	drwxr-xr-x	211.8 kB
▼ lib	0	-	-	12	drwxrwxrwx	2.5 MB
▼ x86_64-linux-gnu	0	-	-	12	drwxrwxrwx	2.5 MB
ld-2.27.so	0	-	-	-	-rwxr-xr-x	179.2 kB
ld-linux-x86-64.so.2 → ld-2.27.so	0	-	-	-	-Lrwxrwxrwx	0
libc-2.27.so	0	-	-	-	-rwxr-xr-x	2.0 MB
libc.so.6 → libc-2.27.so	0	-	-	-	-Lrwxrwxrwx	0
libnss_dns-2.27.so	0	-	-	-	-rwxr--r--	26.9 kB
libnss_dns.so.2 → libnss_dns-2.27.so	0	-	-	-	-Lrwxrwxrwx	0
libnss_files-2.27.so	0	-	-	-	-rwxr--r--	47.6 kB
libnss_files.so.2 → libnss_files-2.27.so	0	-	-	-	-Lrwxrwxrwx	0
libpthread-2.27.so	0	-	-	-	-rwxr-xr-x	145.0 kB
libpthread.so.0 → libpthread-2.27.so	0	-	-	-	-Lrwxrwxrwx	0
libresolv-2.27.so	0	-	-	-	-rwxr--r--	97.1 kB
libresolv.so.2 → libresolv-2.27.so	0	-	-	-	-Lrwxrwxrwx	0
▶ lib64	0	-	-	1	drwxrwxrwx	0
▶ root	0	-	-	1	drwx-----	6.9 MB

## slim/go/app - 9.9MB

- App with deps
- Binary and non-binary deps included automatically
- **/root/go/src/api/api** - 6.9 MB
- <https://github.com/docker-slim/examples/tree/master/3rdparty/mux-go-api>



```
FROM node:18
```

```
WORKDIR /usr/src/app
```

```
COPY app/package*.json ./
```

```
RUN npm install
```

```
COPY app/ .
```

```
EXPOSE 8080
```

```
CMD [ "node", "server.js" ]
```

**slim app** - 90.9MB

- fat(1GB) -> slim(90MB)
- auto-“scratch”

# Kubernetes Debugging Techniques that Don't Work with Slim Images (and Why)



- **Embedded** debugging tools
- `kubectl cp`
- `kubectl exec`

Requires tar to be there in the container, else it will fail

kubectl cp ↗

Copy files and directories to and from containers

Kubectl exec → Needs a shell to be there in the container bash/sh

Debugging utilities → htop, curl etc

# Minimal Container Images: “kubectl exec” failure demo



```
% export PNAME=`kubectl get pods -l app=kubeday-demo -o jsonpath='{.items[0].metadata.name}'`  
% kubectl exec -it -c app $(PNAME) -- bash
```

```
error: Internal error occurred: error executing command in container: failed to exec in container:  
failed to start exec "f53291c5312b9186c9a67adb1c350da51030c1f44e8f8d2e5d3c9430128f3a4f": OCI runtime  
exec failed: exec failed: unable to start container process: exec: "bash": executable file not found in  
$PATH: unknown
```

# Ephemeral Containers: Background

- Why/what?
- History
- How they are implemented in Kubernetes
- How they are exposed through **kubectl debug**

# Ephemeral containers

```
graph TD; A[Ephemeral containers] --> B[These are the containers to help debug pods where there is no way debug directly using exec etc.]; B --> C[It attaches the container to the same pod and share a few namespaces as well]; C --> D[No readiness, liveness probes  
Resources is disallowed]; D --> E[created using ephemeralcontainers API]; E --> F[You cannot expose PORT];
```

These are the containers to help debug pods where there is no way debug directly using exec etc.

It attaches the container to the same pod and share a few namespaces as well

No readiness, liveness probes  
Resources is disallowed  
created using ephemeralcontainers API

You cannot expose PORT

# Ephemeral containers

```
kubectl debug -it nginx --image busybox -- /bin/sh
```



Adds a new container with busybox image to nginx image



Use same resources as the pods



helps to debug containers without any debugging utilities

```
ephemeralContainers:
- command:
  - /bin/sh
  image: busybox
  imagePullPolicy: Always
  name: debugger-cnvwf
  resources: {}
  stdin: true
  terminationMessagePath: /dev/termination-log
  terminationMessagePolicy: File
  tty: true
```

cgroup

mnt

pid

```
ephemeralContainerStatuses:
- containerID: containerd://5659598ca268827dc60472336958afd5e5ebc88ac563c8aadf4
  1660bf5eae26e
  image: docker.io/library/busybox:latest
  imageID: docker.io/library/busybox@sha256:fcd85228d7a25feb59f101ac3a955d27c80
  df4ad824d65f5757a954831450185
```

# Debugging Capabilities



- **Pod level** debugging
  - **same** pod
  - pod **copy** for advanced debugging
    - pod copy as a replacement for the target pod
    - customizing debugged container properties
- **Node level** debugging

# Simple kubectl debug Demo



```
% export PNAME=`kubectl get pods -l app=kubeday-demo -o jsonpath='{.items[0].metadata.name}'`  
% kubectl debug -it -c debug-sidecar-bbox --image busybox --target app ${PNAME}
```

Targeting container "app". If you don't see processes from this container it may be because the container runtime doesn't support this feature.

If you don't see a command prompt, try pressing enter.

```
/ # ps aux  
PID    USER      TIME  COMMAND  
   1   root         0:00 node server.js  
  13   root         0:00 sh  
  18   root         0:00 ps aux
```

```
/ # echo $$  
13
```

```
/ # ls  
bin  dev  etc  home  proc  root  sys  tmp  usr  var
```

```
/ # ls /proc/1/root  
bin  dev  etc  lib  lib64  proc  run  sys  tmp  usr  var
```

# Ephemeral Containers Without "kubectl debug"

- internals/APIs
- additional capabilities (e.g., setting security context, mounting volumes)
- curl examples

# Gotchas

- can't remove ephemeral containers
- not all container properties are available
- process namespace sharing gotchas
- **kubectrl debug** gotchas (can't set security context or mount volumes)

# Ephemeral Containers: Support

- kubernetes version 1.23(beta), 1.25(GA)
- Ephemeral containers support in different hosted Kubernetes services:
  - Civo Kubernetes
  - GKE (Google)
  - EKS (AWS)
  - AKS (Microsoft/Azure)
  - DOKS (DigitalOcean)



# Debugging Container Images

- Custom Debugging Images:
  - **Netshoot**
  - **KoolKits** by LightRun (jvm, node, python, go)
- Do It Yourself:
  - Dockerfile <- stuff :-)
  - <http://nixery.dev> (do it yourself, but on demand)

# Koolkit/Netshoot/Nixery kubectl debug Demos



```
export PNAME='kubectl get pods -l app=kubeday-demo -o jsonpath='{.items[0].metadata.name}'`  
kubectl debug -it -c debug-sidecar-kk --image lightruncom/koolkits:node --target app ${PNAME}
```

Targeting container "app". If you don't see processes from this container it may be because the container runtime doesn't support this feature.  
If you don't see a command prompt, try pressing enter.

```
root@kubeday-demo-657564c8c9-g5wdw:/usr/local/bin# ps aux  
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND  
root         1  0.0  0.8 597528 51148 ?        Ssl   06:09   0:00 node server.js  
root        38  0.1  0.1   7944  7264 pts/0    Ss    06:51   0:00 /bin/bash  
root       205  0.0  0.0   5904  2844 pts/0    R+    06:52   0:00 ps aux
```



```
export PNAME='kubectl get pods -l app=kubeday-demo -o jsonpath='{.items[0].metadata.name}'`  
kubectl debug -it -c debug-sidecar-netshoot --image nicolaka/netshoot --target app ${PNAME}
```

Targeting container "app". If you don't see processes from this container it may be because the container runtime doesn't support this feature.  
If you don't see a command prompt, try pressing enter.

```
                                dP                                dP                                dP  
                                88                                88                                88  
88d888b. .d8888b. d8888P .d8888b. 88d888b. .d8888b. .d8888b. d8888P  
88' `88 88oooo88 88 Y8oooo. 88' `88 88' `88 88' `88 88  
88 88 88. ... 88 88 88 88 88 88. .88 88. .88 88  
dP dP `88888P' dP `88888P' dP dP `88888P' `88888P' dP
```

Welcome to Netshoot! (github.com/nicolaka/netshoot)

kubeday-demo-657564c8c9-g5wdw ▶ ~ ▶



```
export PNAME='kubectl get pods -l app=kubeday-demo -o jsonpath='{.items[0].metadata.name}'`  
kubectl debug -it -c debug-sidecar-nix --image nixery.dev/shell/which/ls/ps/lproute2/netcat-gnu/tshark/tcpdump/curl/jq/nodejs --target app ${PNAME}
```

Targeting container "app". If you don't see processes from this container it may be because the container runtime doesn't support this feature.  
If you don't see a command prompt, try pressing enter.

```
bash-5.1# ps aux  
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND  
0           1  0.0  0.8 597528 51148 ?        Ssl   06:09   0:00 node server.js  
0          265  0.1  0.0   4852  4008 pts/0    Ss    07:02   0:00 bash  
0          271  0.0  0.0   7104  2480 pts/0    R+    07:02   0:00 ps aux
```

```
bash-5.1# ls  
bin dev etc include lib libexec nix nix-support proc sbin share sys tmp usr
```

# Key Takeaways

- **Minimal container images on k8s are ready for mainstream use!!!**
- Ephemeral containers make it possible to debug Scratch, Slim and Distroless app images in production.
- Use tools like Nixery, Netshoot, Koolkits for faster debugging

**Thank You!**

- **Demo** repo:
  - <https://github.com/kubeday-japan/demo>
- Try **ephemeral** containers on Civo!
- Create **minimal container images** with docker-slim or Slim!
- And thanks to **Akihiro Suda** for creating **nerdctl!!!**